
WINE REVIEWS

Gauthier Cler
gauthier.cler@gmail.com

Anas Buyumad
anas.buyumad@epitech.eu

February 2, 2020

Abstract

This paper defines and explains through chronological steps the approach adopted during this data science and visualisation project. We justify our implementation choices, library uses and we interpret obtained results and performances. Finally we suggest areas of improvement on current baseline and relate on dataset applicability.

Contents

1	Introduction	2
2	Approach	2
2.1	Dataset selection	2
2.2	Exploratory data analysis	2
2.3	Data preprocessing	3
2.4	Model selection	5
2.5	Model evaluation	6
2.6	Model tuning	6
3	Wine Recommender	6
4	Conclusion	7

1 Introduction

The goal of the project is to propose visualisations of a dataset and a quantitative analysis.

Project source code, project description and all additional resources are available at

<https://github.com/gauthiercler/wine-reviews>

2 Approach

We define our approach towards this problem in different distinct steps, detailed in the following sections.

2.1 Dataset selection

Among several millions of different datasets of many kinds available online, we decided to go with **wine-reviews**, a popular dataset available on Kaggle, an online community Data Science platform.

<https://www.kaggle.com/zynicide/wine-reviews>

It is about 130,000 rows of data, shared across 13 columns (plus 1 column for id). Each entry represents a detailed review on a wine from an user of www.wineenthusiast.com website.

Dataset Description		
Column name	Data type	Missing values
country	text	63
description	text	0
designation	text	37465
points	numeric	0
price	numeric	8996
province	text	63
region_1	text	21247
region_2	text	79460
taster_name	text	26244
taster_twitter_handle	text	31213
title	text	0
variety	text	1
winery	text	0

One of the reasons for choosing this dataset is the large context diversity on the different columns. Furthermore, the description field is quite long.

2.2 Exploratory data analysis

After describing dataset structure in the last section, let's perform an exploratory analysis on columns and try to find possible correlations between fields. In this dataset, only two columns are of type numeric. Using matplotlib and seaborn visualisation libraries, we can plot distribution. *Points* can be defined as discrete data $\mathbb{D} = \{80, 81, 82, \dots, 100\}$, while *Price* is continuous. We can see that *Price* frequency is roughly centered on $x \in [0, 100]$ (Figure 1b).




Figure 1: Numerical data fields distribution

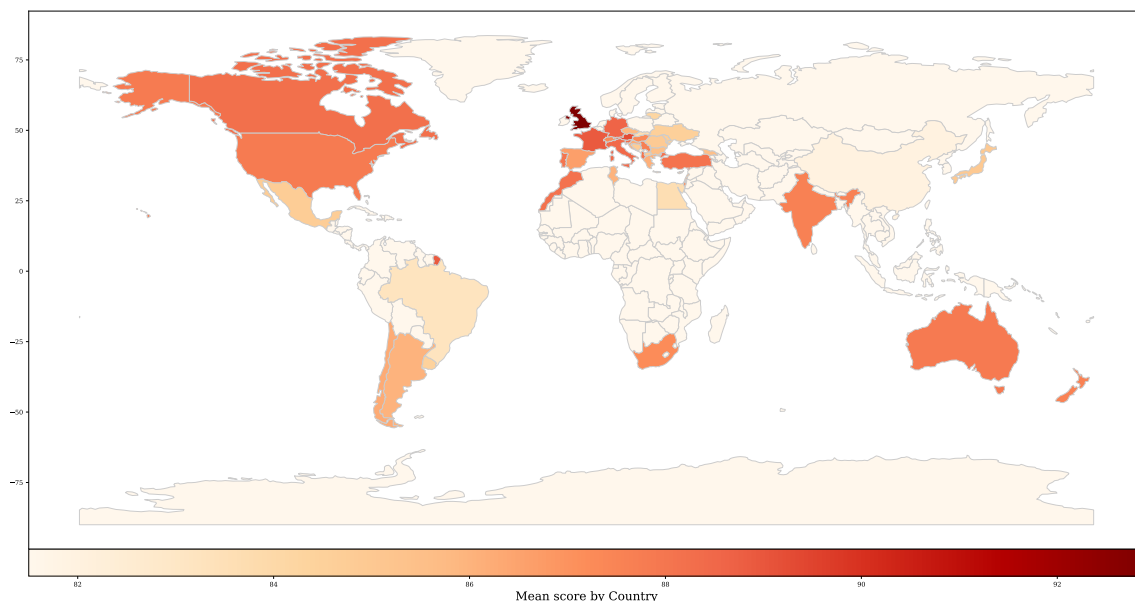


Figure 2: Mean Score by country map

By analyzing the distribution of our data, we can observe that the points of the wines follow a normal distribution (figure 1a) and that their price follow a log-normal distribution (figure 1b). By plotting the evolution of a wine's price over its attributed points (figure 1c), we can recognize an exponential distribution, indicating a correlation between a wine's quality and price. We can thus validate the trivial assumption that the better the wine, the pricier it is.

For non-numerical data, it is not mathematically possible to detect correlations. We can however apply wine-related general knowledge in order to get the most out of it.

For example, we know for a fact that France is the top wine region in the world, while red wines tend to be more expensive than white wines (source : <https://www.vivino.com/wine-news/how-much-does-a-good-bottle-of-wine-cost>, <https://winefolly.com/lifestyle/top-wine-regions-of-the-world/>). Though this kind of information alone is not very pertinent, when coupled with wine-specific characteristics contained in the description field we can gain valuable insight on the wine we are trying to analyze. One such insight would be a wine qualified as having great quality tannins, meaning a higher quality wine, which in turn means a higher price especially if it is a French red wine.

Thus, it is in our interest to process information such as origin, variety or description because an enormous amount of correlations can be found and deduced.

2.3 Data preprocessing

TODO : Layout, italic etc

Before training and testing a model on our dataset, we must pre-process its data according to our analysis in order to get the best predictions possible.

Firstly, we randomly sort our dataset and remove duplicate entries.

```
X = wines.sample(frac=1).reset_index(drop=True)
X.drop_duplicates(inplace=True)
```

We then choose to drop certain columns for different reasons:

- Designation is dropped because 29% of its entries are null values, with almost 40.000 unique values. Knowing the vineyard where the grapes were used for the wine is too specific and would not give us any useful insight especially when contemplating the training time cost along with the sheer number of potential features.
- Winery is dropped for the same reasons as Designation, 16.757 unique values is too specific, and knowing in which winery the vineyard is does not add much to quality prediction.
- Region_1 is dropped because 16% of its entries are null values and 1.228 unique values giving the place inside a province or state where the wine was grown is too costly all the while giving very low benefits to prediction accuracy. We already possess the Province feature giving us a broader and more accurate insight.
- Region_2 is dropped for the same reasons as Region_1, being just a more precise version of it, boasting 61% of null values.
- Taster_name and Taster_twitter_handle are dropped because knowing who rated the wine is useless for our case.
- Title is dropped for being too specific with an expected 118.840 values, meaning one unique value for every non-duplicated wine.

Finally, we simply remove any wine that possesses one or more feature with a null value.

```
X.drop(['designation', 'winery', 'region_1', 'region_2', 'taster_name',
'taster_twitter_handle', 'title'], inplace=True, axis=1)
X.dropna(axis=0, inplace=True)
```

In order to gain worthwhile information from our wine's Description, we first need to reduce its complexity. We use a count vectorizer to sort the words (mainly adjectives) by their amount of occurrences throughout the dataset, allowing us to extract and keep meaningful key-words with the best potential of characterizing our wines (1000 words in our case). English stop words are also used to reject common words.

We end up with a two dimensional array, representing the weight of each key-word according to each wine, that we merge to the original dataset while dropping the description column.

```
from sklearn.feature_extraction.text import CountVectorizer
from scipy.sparse import hstack

vectorizer = CountVectorizer(stop_words='english')
result = vectorizer.fit_transform(X['description'])
feature_array = vectorizer.get_feature_names()
top_words = sorted(list(zip(vectorizer.get_feature_names(),
                             result.sum(0).getA1()),
                     key=lambda x: x[1], reverse=True)[:1000])

words_indexes = [feature_array.index(x[0]) for x in top_words]
result = result[:, indexes]
X_merged = hstack((X, result))
```

For the rest of the data (Country, Province and Variety), we use a one-hot encoding method in order to numerically represent the values as binary. Meaning, in the case of the Country for example, 44 different features with only one having a value set to 1, stored as a sparse matrix in order to save space. This effectively allows us to tell our model where a wine comes from or its variety, getting rid of the nominal nature of this information.

Before doing so, we also isolate the Points from the dataset.

```
from sklearn.preprocessing import OneHotEncoder

y = X['points']
X.drop('points', axis=1, inplace=True)
```

```
encoder = OneHotEncoder()
X = encoder.fit_transform(X)
```

Finally, we split our processed data into two group, train and test.

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, shuffle=True)
```

We define 90% as **train** set and 10% as **test** set. The goal here is to train the model on the train test, so it can generalize and create an intuition on the data. We can then validate using the test set to evaluate model understanding and reliability.

2.4 Model selection

When it comes to model selection, it's not always easy to find an estimator exactly suited for a given problem and data. We can group estimators in 4 categories regarding data problems (see Figure 3). In our case, we want to predict wine quality (column points) from other features. As we saw in EDA section, points are distributed from 80 to 100. We are dealing with a regression problem. Note that because we are working with discrete distribution, we can also define this problem as a classification problem with 20 categories. We can use scikit-learn flowchart to help us to find the most appropriate estimator for our data and goals.

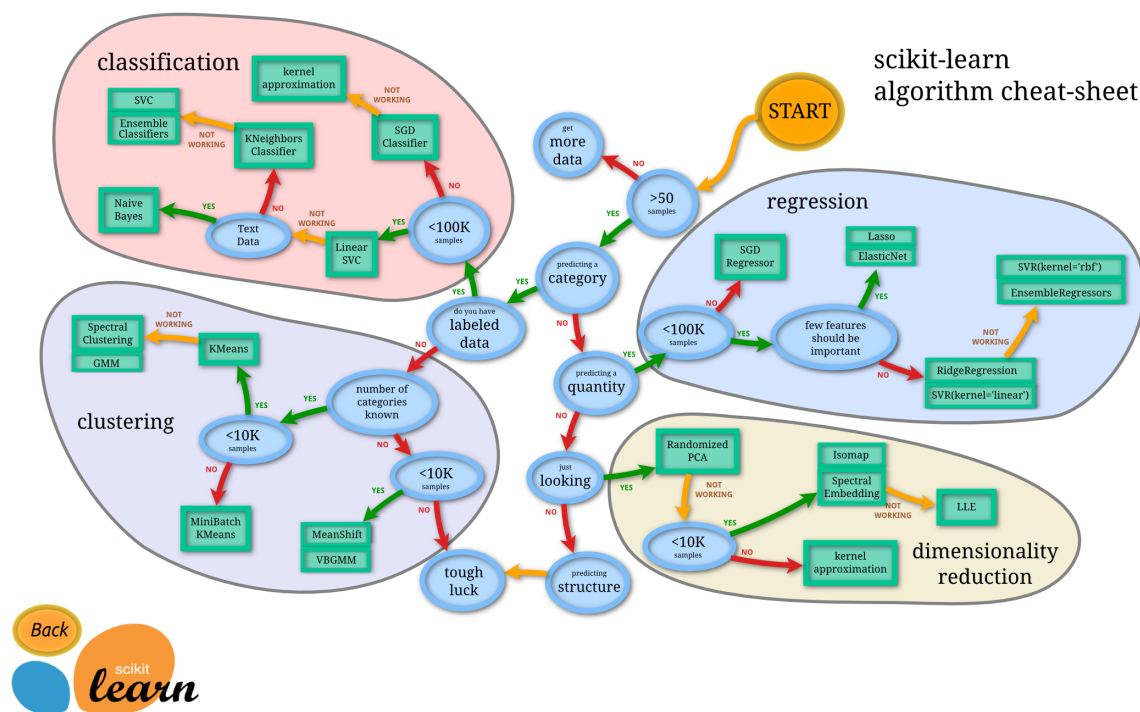


Figure 3: scikit-learn estimator by problem flowchart

As our cleaned data is not too big, variables are independent and correlations shows that only some features represent y, we can try to apply a linear Regression like SGD, Lasso or ElasticNet. However, After empirical tests, we decide to move to ExtraTreeRegressor, as random split truly improve training speed without losing any accuracy.

```
from sklearn.ensemble import ExtraTreeRegressor
```

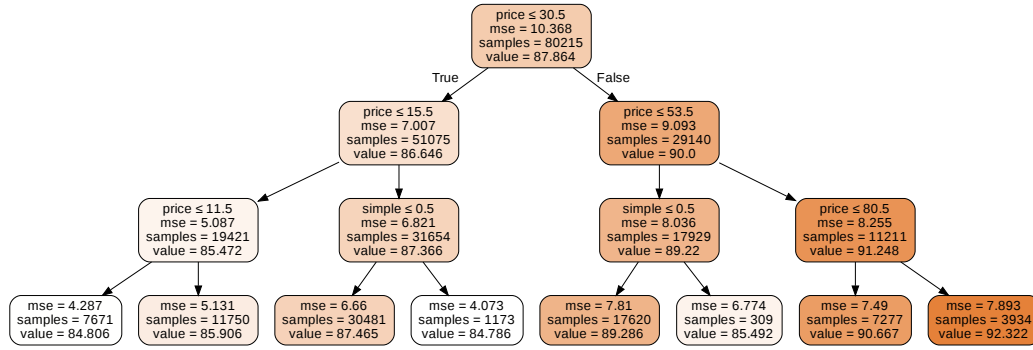


Figure 4: Subset of a decision tree

```

clf = ExtraTreeRegressor()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

```

The way we are evaluating model is defined in the next section.

2.5 Model evaluation

As we are dealing with regression, we decided to use **Mean Absolute Error** (MAE) metric to evaluate how well is performing RandomForest on the data.

MAE is defined by:

$$MAE(y, \hat{y}) = \frac{1}{n_{samples}} \sum_{i=0}^{n_{samples}-1} |y_i - \hat{y}_i|. \quad (1)$$

Using chosen estimator with default parameters, we achieve MAE of **1.712**. We can easily calculate MAE with scikit-learn:

```

from sklearn.metrics import mean_absolute_error

mean_absolute_error(y_true, y_pred)

```

$$R^2(y, \hat{y}) = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \quad (2)$$

To reduce loss, we will tune hyper parameters of current models to find the best combination.

2.6 Model tuning

Speak about max features: sqrt / SelectKBest

3 Wine Recommender

Based on dataset analysis and wine review description, we made a website to recommend wines. By defining your maximum price, filtering by Country / wine variety, and by describing your ideal wine with keywords such as "soft", "fruity" or "dry", this tool will be able to recommend you a list of wines that would match your needs.

It is available at:

<http://wine-recommender.gauthierclerc.com/>

4 Conclusion

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.